

The Few Parts That Matter

Why this exists

Companies write a lot of documents. Design docs, requirement docs, planning docs, notes someone dumped two years ago. They pile up fast.

Here's the thing nobody says out loud: inside any one of those documents, most of it doesn't matter to you right now. A small part does. Maybe there's one decision buried on page 12 that, if you miss it, you build the wrong thing or break something that was working. The rest is background you can skip.

And which part matters keeps changing. If you're working on payments today, one part is important. If you're working on login tomorrow, a completely different part matters — and the payments stuff is now just noise.

Now imagine not one document, but a few hundred of them, written by people who've since left. Somewhere in that pile are the handful of parts that matter for the exact thing you're about to do. Finding them is slow, and mostly you don't even know they exist, so you don't look.

That's the real problem. Not "store the documents." Not "link them together." **Show me the few parts that matter for what I'm doing right now — and let me trust them, dig deeper when I want, and keep reading as if nothing clever were happening.**

What tools do today

Some tools store your notes and let you link one to another. Some track your tasks. They're fine. But they all work at the level of the whole document — you open a document, you read a document, you search for a document.

None of them help with the actual pain: you don't want the document. You want the few parts inside it that matter to you today — and you want to know you can rely on them.

The question we want you to sit with: everyone stops at "here's your document." Should you?

A quick word on "a part"

We keep saying "a part" on purpose, because it isn't always one line.

Sometimes the important thing is a single sentence. Sometimes it's a whole paragraph. And sometimes it isn't fixed text at all — it's a live chart whose numbers are pulled from a system

somewhere and change over time. Same idea, very different shapes. Your design has to handle all of them.

What to build

A simple wireframe/web page/app for moving through a big, messy pile of documents — where the thing you see and pick from is **the important parts inside them, not the whole files**.

Assume the "which part matters" question is already answered (more on that below). Your job is everything the person sees and touches. This is where the real work is, so here are the questions we care most about:

Showing a part, and where it's from. When you show an important part, how do you also show where it came from — which document, which section — without cluttering the screen? Enough that the person believes it and can go check.

Handling different shapes cleanly. A sentence, a full paragraph, and a live chart pulling fresh numbers — how does your design show all three the same calm way, without the layout falling apart or each one turning into a special case?

Letting them go deeper for the full story. Sometimes the part isn't enough — they want the whole context around it. How do they jump into the full document, land right on that part in its original place, look around, and get back to where they were without losing their spot?

Helping them decide whether to trust it. Who wrote this? When? Is it still true, or from an abandoned doc that's three years stale? What's it based on? Show enough that someone can decide in about a second whether to lean on it.

Make it clear who wrote and owns each part. A single document often carries parts from different hands — different people, teams, or sources. As someone reads or browses, it should always be clear whose words these are and where a part came from, without them having to go digging. Make that authorship visible while reading — there when it matters, out of the way when it doesn't. And when a part on the page came from somewhere else, let the person decide what to do with it: **take it in as their own, keep it as a pointer back to its source, or drop it**. That choice should feel quick and clear.

Reading straight through should feel like a plain document. This is the heart of the assignment, and it's not really about speed. It's about feel and space. If someone reads a document top to bottom, it should feel as clean and as space-efficient as reading a plain text file, a markdown file, or a PDF — the same calm, the same tight use of the page, nothing wasted. All your clever surfacing, trust signals, authorship cues, and controls should step back and get out of the way when the person is simply reading. The moment the page feels like a busy app instead of a document you're reading, you've lost the thing that made documents good in the first place.

That's the real tension: powerful when you need it, and completely calm — as calm as plain text — when you don't.

You won't do all of this well in the time you have. **Pick the two or three that you think matter most, go deep, and tell us why you set the rest aside.** That choice is part of what we're looking at.

One thing you can skip: you don't have to build the smart part that figures out which part is important. Fake it — tag things by hand, match on keywords, whatever's easy. We're not testing that. We're testing how the thing looks and feels to use. Make the smart part pretend; make the experience real.

We are also sharing - a realistic, fully fictional company knowledge base, built as the working material to support this assignment. Use/change it or make your own — just make it feel real, not twenty tidy notes. A tidy pile hides the whole problem.

The one rule

Use AI for as much of this as you want. 90% AI is fine, honestly expected.

We just want to see the 10% that's you: the idea the AI wouldn't have had, the moment you didn't accept its answer and did something better, the small thing you cared about that nobody asked you to. Point those out as you go.

What we're actually testing

- When you know a part matters, do you show it — and where it came from — in a way that's clear and trustworthy, not cluttered?
- Can you show a sentence, a paragraph, and a live chart all the same calm way?
- Can someone move from a single part, out to the full document for context, and back, without ever feeling lost?
- Can a person tell in a glance whether to trust something — who, when, still-true-or-not?
- While reading, is it clear who wrote and owns each part — and can the person easily take in, link to, or drop a part that came from another source?
- When someone just reads a document end to end, does it feel as clean and space-efficient as a plain text or PDF file — or like a busy app?
- Can you point at something and say "I chose not to build this, on purpose, because..." and be right?
- Does building this sound fun, or like a chore?

Logistics

- **A wireframe is fine; a working app is what we'd love to see.** If you build one, a small JavaScript app backed by made-up, hard-coded data & files is perfect. No real backend, no real run time document fetches — fake the data, documents, build the experience.
- A short note — a few paragraphs — on what you decided, and what you left out on purpose. Short and sharp beats long.

- Do a few things well rather than many things halfway.
- Don't over-polish. A rough edge you can explain beats a shiny one you can't.

Build the thing that hands someone the few parts that matter, lets them trust those parts and dig in when they want — and, when they just want to read, feels exactly as calm and clean as a plain text file.